

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3
4  def init_two_layer_model(input_size, hidden_size, output_size):
5      """
6      Initialize the weights and biases for a two-layer fully connected neural
7      network. The net has an input dimension of D, a hidden layer dimension of H,
8      and performs classification over C classes. Weights are initialized to small
9      random values and biases are initialized to zero.
10
11      Inputs:
12      - input_size: The dimension D of the input data
13      - hidden_size: The number of neurons H in the hidden layer
14      - output_size: The number of classes C
15
16      Returns:
17      A dictionary mapping parameter names to arrays of parameter values. It has
18      the following keys:
19      - W1: First layer weights; has shape (D, H)
20      - b1: First layer biases; has shape (H,)
21      - W2: Second layer weights; has shape (H, C)
22      - b2: Second layer biases; has shape (C,)
23      """
24      # initialize a model
25      model = {}
26      model['W1'] = 0.00001 * np.random.randn(input_size, hidden_size)
27      model['b1'] = np.zeros(hidden_size)
28      model['W2'] = 0.00001 * np.random.randn(hidden_size, output_size)
29      model['b2'] = np.zeros(output_size)
30      return model
31
32  def two_layer_net(X, model, y=None, reg=0.0):
33      """
34      Compute the loss and gradients for a two layer fully connected neural network.
35      The net has an input dimension of D, a hidden layer dimension of H, and
36      performs classification over C classes. We use a softmax loss function and L2
37      regularization the the weight matrices. The two layer net should use a ReLU
38      nonlinearity after the first affine layer.
39
40      The two layer net has the following architecture:
41
42      input - fully connected layer - ReLU - fully connected layer - softmax
43
44      The outputs of the second fully-connected layer are the scores for each
45      class.
46
47      Inputs:
48      - X: Input data of shape (N, D). Each X[i] is a training sample.
49      - model: Dictionary mapping parameter names to arrays of parameter values.
50        It should contain the following:
51      - W1: First layer weights; has shape (D, H)
52      - b1: First layer biases; has shape (H,)
53      - W2: Second layer weights; has shape (H, C)
54      - b2: Second layer biases; has shape (C,)
55      - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
56        an integer in the range 0 <= y[i] < C. This parameter is optional; if it
57        is not passed then we only return scores, and if it is passed then we
58        instead return the loss and gradients.
59      - reg: Regularization strength.
60
61      Returns:
62      If y not is passed, return a matrix scores of shape (N, C) where scores[i, c]
63      is the score for class c on input X[i].
64
65      If y is not passed, instead return a tuple of:
66      - loss: Loss (data loss and regularization loss) for this batch of training
67        samples.

```

```

68 - grads: Dictionary mapping parameter names to gradients of those parameters
69 with respect to the loss function. This should have the same keys as model.
70 """
71
72 # unpack variables from the model dictionary
73 W1,b1,W2,b2 = model['W1'], model['b1'], model['W2'], model['b2']
74 N, D = X.shape
75
76 # compute the forward pass
77 scores = None
78 #####
79 # TODO: Perform the forward pass, computing the class scores for the input. #
80 # Store the result in the scores variable, which should be an array of #
81 # shape (N, C). #
82 #####
83
84 # evaluate class scores with a 2-layer Neural Network
85 hidden_layer = np.maximum(0, np.dot(X, W1) + b1) # note, ReLU activation
86 scores = np.dot(hidden_layer, W2) + b2
87
88 #####
89 #                                     END OF YOUR CODE                                     #
90 #####
91
92 # If the targets are not given then jump out, we're done
93 if y is None:
94     return scores
95
96 # compute the loss
97 loss = None
98 #####
99 # TODO: Finish the forward pass, and compute the loss. This should include #
100 # both the data loss and L2 regularization for W1 and W2. Store the result #
101 # in the variable loss, which should be a scalar. Use the Softmax #
102 # classifier loss. So that your results match ours, multiply the #
103 # regularization loss by 0.5 #
104 #####
105
106 num_examples = X.shape[0]
107 # get unnormalized probabilities
108 exp_scores = np.exp(scores)
109 # normalize them for each example
110 probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
111 correct_logprobs = -np.log(probs[range(num_examples),y])
112 # compute the loss: average cross-entropy loss and regularization
113 data_loss = np.sum(correct_logprobs)/num_examples
114 reg_loss = 0.5*reg*np.sum(W1*W1) + 0.5*reg*np.sum(W2*W2)
115 loss = data_loss + reg_loss
116
117 #####
118 #                                     END OF YOUR CODE                                     #
119 #####
120
121 # compute the gradients
122 grads = {}
123 #####
124 # TODO: Compute the backward pass, computing the derivatives of the weights #
125 # and biases. Store the results in the grads dictionary. For example, #
126 # grads['W1'] should store the gradient on W1, and be a matrix of same size #
127 #####
128
129 # compute the gradient on scores
130 dscores = probs
131 dscores[range(N),y] -= 1
132 dscores /= N
133
134 # W2 and b2

```

```
135 grads['W2'] = np.dot(hidden_layer.T, dscores)
136 grads['b2'] = np.sum(dscores, axis=0)
137 # next backprop into hidden layer
138 dhidden = np.dot(dscores, W2.T)
139 # backprop the ReLU non-linearity
140 dhidden[hidden_layer <= 0] = 0
141 # finally into W,b
142 grads['W1'] = np.dot(X.T, dhidden)
143 grads['b1'] = np.sum(dhidden, axis=0)
144
145 # add regularization gradient contribution
146 grads['W2'] += reg * W2
147 grads['W1'] += reg * W1
148
149 #####
150 #                                END OF YOUR CODE                                #
151 #####
152
153 return loss, grads
154
155
```